



One platform, three homework routes

Scheduled agents review evidence; shared Python code manages collection, records and delivery.



From message to verified reply

One run, one lock, one reviewed plan. Student replies use verified data already held by the workflow.



Stop and report: failed / partial collection, stale plan or uncertain delivery. Keep unclear identity or evidence pending.

Replies informed by saved data

Verified preferred name

Use an English greeting; otherwise "Hi,".

Assignment + submission history

Name the task and acknowledge added working or a sentence rewrite when recorded.

QQ grading results

Saved score + one improvement sentence + the assignment model-answer image.

Where repeated work is reduced

Compact packets and context

Review pending evidence; avoid replaying old inboxes, setup discussions and logs.

Cached OCR + bounded reads

Reuse unchanged transcriptions; QQ overlaps independent history reads.

Observed S3.6 compaction

498k to 45.4k context tokens: 90.9% less context, not a cost or speed guarantee.

Shutdown and morning restart

Computer off = checks paused

Local agents cannot execute while the machine and their apps are closed.

Resume the same conversations

Reopen Codex and Kimi Code; restore QQ transport and sign in if requested.

Next successful scan

Use saved state to collect outstanding work. Confirm the first run completes.